

1. Basic Structure of a Relation

- **Relation:** In the relational model, a relation is represented as a table. The table consists of:
 - **Attributes** (columns): These define the type of data that can be stored. Each attribute has a name and a domain (set of allowed values).
 - **Tuples** (rows): Each row represents a single data record. The number of rows is called the **cardinality** of the relation, and the number of columns is called the **degree** of the relation.
- **Example:** In a table called **Students**, the attributes could be **StudentID**, **Name**, **Age**, and the tuples represent individual students:

StudentID	Name	Age
1	Ritesh	20
2	Priya	19
3	Aman	22

2. Domain

- A **domain** is the set of all possible values an attribute can take. For example:
 - The domain of the **Age** attribute might be integers between 0 and 120.
 - The domain of **Name** might be strings up to 50 characters.
- Every attribute in a relation is associated with a domain to ensure consistency and avoid invalid data.

3. Keys

- **Primary Key:** This is a unique identifier for each tuple in a relation. It ensures that no two rows have the same value for the primary key.
 - **Example:** In the **Students** table, **StudentID** is the primary key, meaning every student must have a unique **StudentID**. No two rows can have the same ID, ensuring uniqueness.

- **Candidate Key:** A candidate key is any set of attributes that could serve as a primary key. Out of all candidate keys, one is chosen as the primary key.
 - Example: **StudentID** and maybe an email address could be candidate keys because both uniquely identify a student.
- **Foreign Key:** This key is used to create a relationship between two tables. A foreign key in one table refers to the primary key in another table. It helps establish relationships and maintain referential integrity.
 - Example: A **Grades** table might have a **StudentID** column as a foreign key, linking it to the primary key **StudentID** in the **Students** table.

4. Relational Database Constraints

These constraints enforce rules on the data to maintain accuracy and consistency:

- **Domain Constraint:** This ensures that values for an attribute must come from the specified domain.
 - Example: If the domain of the **Age** attribute is restricted to integers from 0 to 120, entering **200** as an age would violate the domain constraint.
- **Key Constraint:** This constraint ensures that no two tuples (rows) have the same value for the primary key.
 - Example: In the **Students** table, no two students can have the same **StudentID**.
- **Entity Integrity:** This rule ensures that no primary key can have a **NULL** value, which means every tuple must have a valid, non-null identifier.
 - Example: You cannot have a student record without a **StudentID** in the **Students** table.
- **Referential Integrity:** This ensures that a foreign key must match an existing primary key value or be **NULL**. It enforces consistency between related tables.
 - Example: In the **Grades** table, if a **StudentID** references the **Students** table, every **StudentID** in **Grades** must exist in the **Students** table. If you try to insert a **StudentID** in **Grades** that doesn't exist in **Students**, it will violate referential integrity.